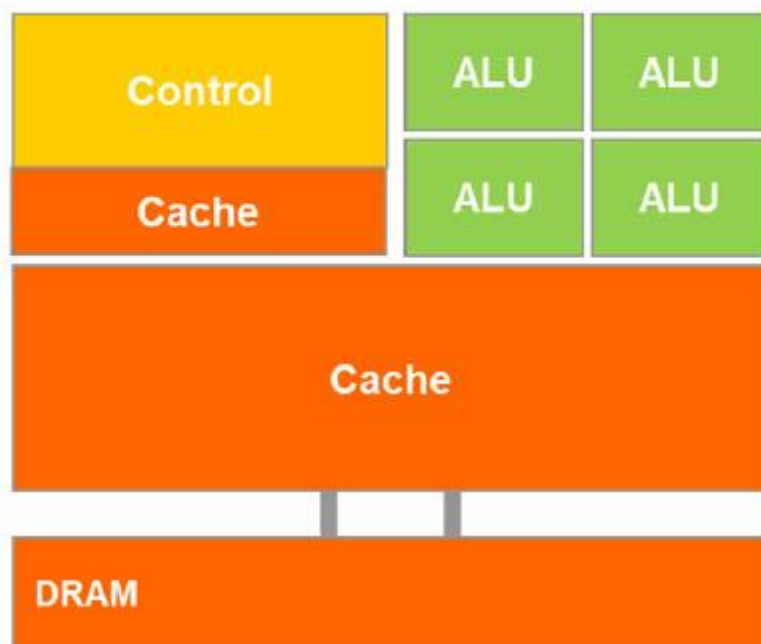


Parallel Programming


- Parallel programming is like getting help from friends to clean up a big mess. Instead of one person doing all the work, everyone works together at the same time. Each person cleans a small part, and when everyone finishes, the whole place is clean faster. This way, the job is done quicker because many people are helping at once.



CPU

(latency-oriented design)

(A)




GPU

(throughput-oriented design)

(B)

FIGURE 1.1

CPUs and GPUs have fundamentally different design philosophies: (A) CPU design is latency oriented; (B) GPU design is throughput-oriented.



```
01 void vecAdd(float* A, float* B, float* C, int n) {
02     int size = n* sizeof(float);
03     float *d_A *d_B, *d_C;
04
05     // Part 1: Allocate device memory for A, B, and C
06     // Copy A and B to device memory
07     ...
08
09     // Part 2: Call kernel - to launch a grid of threads
10     // to perform the actual vector addition
11     ...
12
13     // Part 3: Copy C from the device memory
14     // Free device vectors
15     ...
16 }
```

FIGURE 2.5

Outline of a revised vecAdd function that moves the work to a device.

```
01 void vecAdd(float* A_h, float* B_h, float* C_h, int n) {
02     int size = n * sizeof(float);
03     float *A_d, *B_d, *C_d;
04
05     cudaMalloc((void **) &A_d, size);
06     cudaMalloc((void **) &B_d, size);
07     cudaMalloc((void **) &C_d, size);
08
09     cudaMemcpy(A_d, A_h, size, cudaMemcpyHostToDevice);
10     cudaMemcpy(B_d, B_h, size, cudaMemcpyHostToDevice);
11
12     // kernel invocation code - to be shown later
13     ...
14
15     cudaMemcpy(C_h, C_d, size, cudaMemcpyDeviceToHost);
16
17     cudaFree(A_d);
18     cudaFree(B_d);
19     cudaFree(C_d);
20 }
```

FIGURE 2.8

A more complete version of vecAdd().

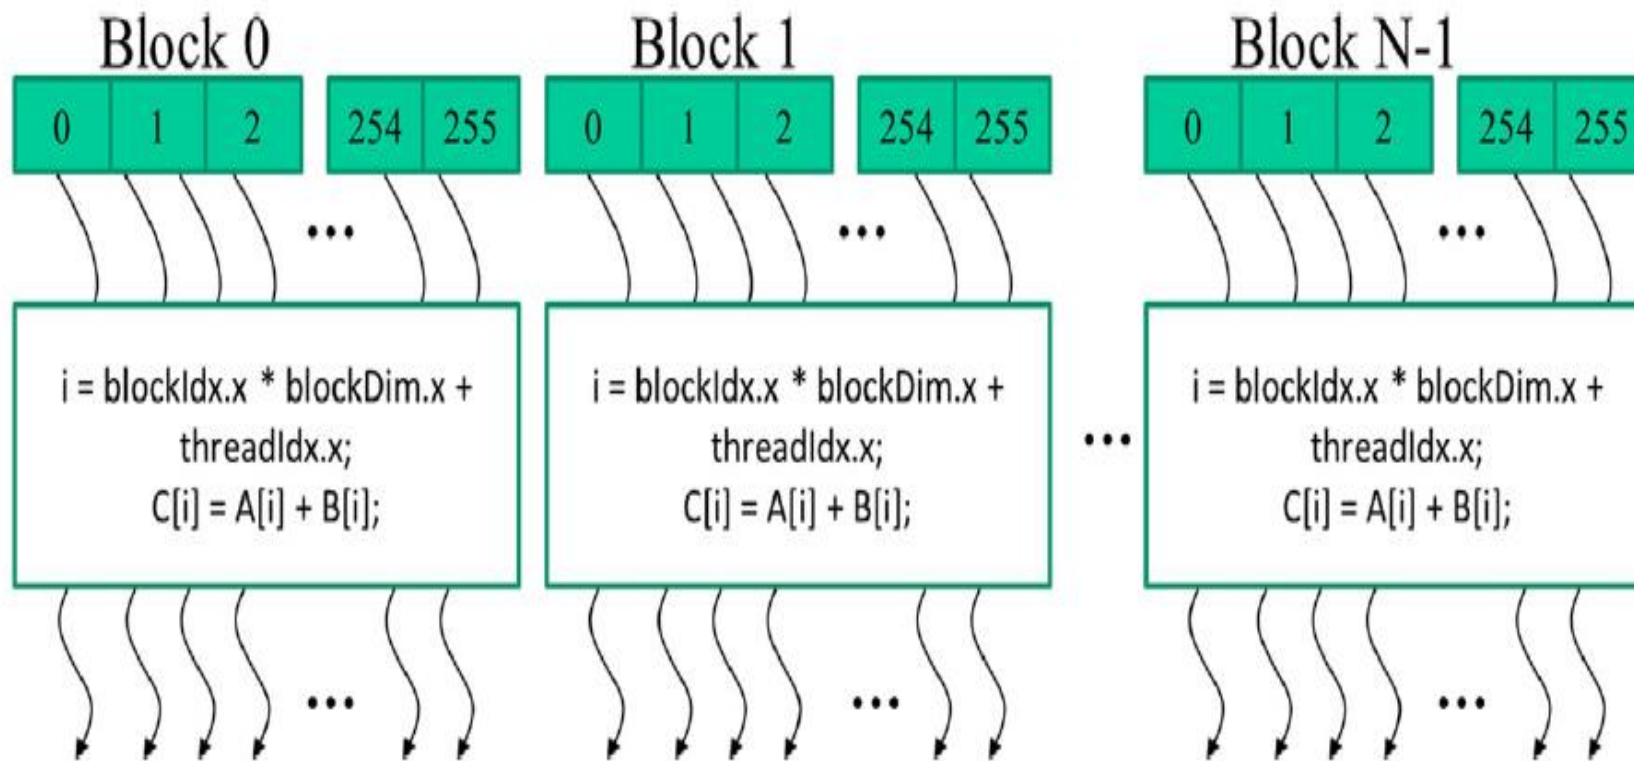


FIGURE 2.9

All threads in a grid execute the same kernel code.

```
01  // Compute vector sum C = A + B
02  // Each thread performs one pair-wise addition
03  __global__
04  void vecAddKernel(float* A, float* B, float* C, int n) {
05      int i = threadIdx.x + blockDim.x * blockIdx.x;
06      if (i < n) {
07          C[i] = A[i] + B[i];
08      }
09  }
```

FIGURE 2.10

A vector addition kernel function.

```
01 void vecAdd(float* A, float* B, float* C, int n) {
02     float *A_d, *B_d, *C_d;
03     int size = n * sizeof(float);
04
05     cudaMalloc((void **) &A_d, size);
06     cudaMalloc((void **) &B_d, size);
07     cudaMalloc((void **) &C_d, size);
08
09     cudaMemcpy(A_d, A, size, cudaMemcpyHostToDevice);
10     cudaMemcpy(B_d, B, size, cudaMemcpyHostToDevice);
11
12     vecAddKernel<<<ceil(n/256.0), 256>>>(A_d, B_d, C_d, n);
13
14     cudaMemcpy(C, C_d, size, cudaMemcpyDeviceToHost);
15
16     cudaFree(A_d);
17     cudaFree(B_d);
18     cudaFree(C_d);
19 }
```